

Barely Broken – Testing Highly Structured Inputs

Valentin Huber

valentin.huber@cispa.de

CISPA Helmholtz Center for Information Security

Germany

Abstract

Current test generators are limited in the inputs they produce. This leads to an imbalance in the program parts tested, with significant program parts reaching insufficient attention, particularly with programs expecting inputs that are highly structured. In this work, I propose two measurements to evaluate different approaches and find which parts of programs are under-tested. *Input correctness* provides a level to which an input satisfies the expected structure, measured either against an implementation or specification. *Input interestingness* checks the distance to equivalence class boundaries.

I present the distribution of *input correctness* from an initial evaluation of different approaches on a C compiler, suggesting significant limitations of existing approaches. I further propose three approaches to closing the gap in testing all input parts.

CCS Concepts

- **Do Not Use This Code → Generate the Correct Terms for Your Paper:** *Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.*

Keywords

Fuzzing, Grammar Fuzzing, Language Fuzzing

ACM Reference Format:

Valentin Huber. 2018. Barely Broken – Testing Highly Structured Inputs. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXXX.XXXXXXXX>

1.change title

1 Introduction

2.number programs in fig

Programs processing input typically do so in several distinct steps, with the programs in Figure 1 passing increasingly more of these steps:

- (1) **Lexing:** Unstructured input is split into an unstructured list of tokens. (P1) will be rejected by the lexer of a C compiler.

Unpublished working draft. Not for distribution.
Permission to make digital or hard copies of all or part of this work for personal or
internal use is granted, provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation
on the first page. Copyrights for components of this work owned by others than the
author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or
republis, to post on servers or to redistribute to lists, requires prior specific permission
and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXXX.XXXXXXXX>

2025-11-15 10:04. Page 1 of 1–4.

(P1): VBq -"7 .6arK w ;zu%6\$K>%OTV" ryD*

(P2): && float += % " (; == [

(P3): int main() { return i; }

(P4): int main() { return 0; }

Figure 1: C programs that are lexically (P1), syntactically (P2) and semantically (P3) invalid and fully valid (P4).

- (2) **Syntactic Parsing and Checks:** If the input is syntactically valid, it is parsed into structures. While (P2) consists of exclusively valid tokens, they do not appear in the correct structure to be parsed into an abstract syntax tree and are therefore rejected by the parser of a compiler.
- (3) **Semantic Checks:** These parsed structures are tested on their semantic meaning. (P3) can be parsed, but is still not valid – the variable i is not defined.
- (4) **Business Logic:** Finally, the input is processed by a program’s inner logic, e.g. translated into an executable.

My observation is that test generators generally produce inputs that test only a subset of the steps above, depending on their internal model of the input structure. Some of my previous work [10, 11] has explored this by building purpose-built fuzzers that ensure (partial) correctness of inputs.

Imagine the input space of a program expecting highly structured input as a bush, with bugs represented by berries spread throughout, as shown in Figure 2.

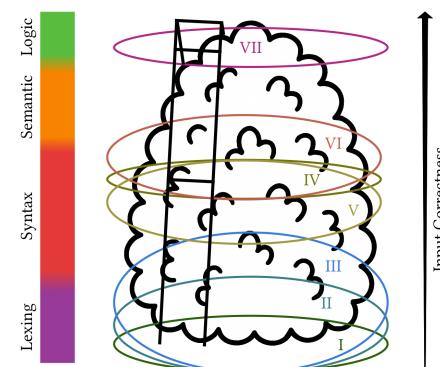


Figure 2: Bush representing the input space of a program, where increasing height represents increasing correctness of a certain input

Purely random input generation, like the first fuzzers [13] is unlikely to test anything but the lexing stage, it is therefore similar to reaching only the very bottom of the bush (I).

In a next step, coverage-guided fuzzing was proposed, which is able to incrementally find inputs that reach deeper into the program with random mutation [18], and thus reaches further up the tree (II). These approaches were extended with instrumentation and logic that allow them to produce increasingly correct inputs [6, 9], thus extending their reach once more (III).

By manually providing the test generator with a model of the target input structure, an analyst can help a fuzzer produce increasingly correct inputs. This is the equivalent of helping a test generator climb on the first step of a ladder.

One widely used approach is to provide a test generator with a context-free grammar, based on which it can produce inputs that will by construction not be rejected by the parser and thus reach further into the program (IV). From there, they can produce inputs that validate their inner model [4, 7] to reach further down (V), or attempt to pass additional semantic checks [8] (VI).

Recent advancements provide a generic model for input testers to receive a complete model of the PUT input structure [19], which allows testing the business logic (VII).

3.Talk about symbex and target-specific fuzzers?

To my knowledge, there is no systematic evaluation of representatives of these approaches – we do not know the extent of their reach. Following this, we do not know what parts of targets remain untested, even when multiple approaches are combined.

Based on this insight, I present the following contributions:

- (1) I propose two axes on which to compare the distribution of inputs from different test generators: input *correctness*, and input *interestingness* (Section 2)
- (2) I provide initial experiments and their results, which suggest that existing input generators are incapable of testing all program logic (Section 3)
- (3) I propose three approaches to creating input generators that fill these holes (Section 4)

2 Evaluating Test Input Quality

Evaluating test generators can be done in multiple, orthogonal ways. The most common measurement is the code coverage reached when executing a certain input on an instrumented implementation of the PUT. While code coverage is comparatively easy to measure due to the availability of instrumentation passes in popular compilers [1, 2]^{4.citation format of these} and binary-only fuzzing tools [12], it has two fundamental limitations: (1) It assigns equal weight to all branches and (2) assumes any value covering a new branch is equally interesting.

2.1 Input Correctness

Limitation (1) of using code coverage in evaluating a fuzzer's performance means that fuzzers that reach far into the PUT but only provide valid inputs may receive the same score as a fuzzer that tests all error paths in the initial input validation steps of a program. Alternatively, I propose evaluating fuzzers for *input correctness*. The correctness of an input with regards to a PUT can be evaluated in two ways:

(1) First, one can test inputs on an implementation of a PUT, measuring the ratio of accepted to rejected inputs at each stage of the input processing pipeline. This can further be refined by annotating or instrumenting the code to record the step at which a certain input is rejected. This can be achieved with one of the following:

- (a) Manual annotations of steps along the path through the program, similar to [5, 15]
- (b) Selective coverage instrumentation of entire program parts as a very coarse measurement [14]
- (c) I further propose the following heuristic: Based on a diverse set of (fuzzer-produced) inputs, *instrument the edges executed by all valid inputs*.

(2) Alternatively, inputs can be evaluated against an abstract specification of the language expected by a PUT. Such specifications are usually given in natural language. Recent works have explored how to express such specification in a structured way, as a combination of a context-free grammar and additional constraints over nodes in this grammar [19]. Evaluating the correctness against such a specification may be done in different steps:

- (a) For inputs that cannot be parsed with the provided grammar, existing literature provides algorithm to measure how close to valid the input is. [3, 16]
- (b) For inputs that are parsed, the ratio of fulfilled to violated constraints can be used to assess the correctness of an input.

2.2 Input Interestingness

Not all inputs are equally interesting, even if they are all fully correct, or are processed by the exact same instructions (as described in Limitation (2)). Assume a constraint that checks whether a person is an adult. According to the principle of equivalence classes, mutating their age to random values, say from 27 to 43, is unlikely to trigger a new bug. Changing it to values near the boundaries of equivalence classes like 18 or 19 is a more promising strategy to discover bugs in edge cases, such as off-by-one errors. Alternatively, one can mutate an input to a generally interesting value, like 0, -1, or 2^{16} [18]^{5.AFL or AFL++?}.

Previously, it was non-trivial to evaluate how close an input part is to such a constraint boundary. With the availability of formal and machine-readable specifications, such as described above, this now becomes possible. An input (part) is more interesting the closer to the boundaries provided by the specification it is. This can be applied to semantic constraints, but also syntactic structure: If a node can be repeated an arbitrary number of times, testing 0 repetitions, 1 repetition and a large number of repetitions is more interesting. This can be reinterpreted as syntactic boundary testing.

An alternative approach to incentivize a fuzzer to produce more interesting inputs is to instrument comparison instructions on parts of the input in a way that reward the fuzzer for inputs that achieve smaller differences between the two values that are compared. This approach is limited by the ability of the instrumentation of the fuzzer to distinguish between comparisons that represent closeness to equivalence class boundaries, as opposed to comparisons that are independent of the input or do not represent boundaries. Absent an

233
234
235
236
**Table 1: Ratio of inputs rejected by program steps and total
237 coverage across approaches in clang. Approaches marked
238 with * receive an initial corpus of valid C programs.**

	Other	Lexing	Syntax	Semantic	Valid	Cov
*	0.000%	0.000%	0.000%	0.000%	1.000%	12772
(IV)	0.086%	0.181%	0.108%	0.623%	0.001%	13074
(IV)*	0.085%	0.177%	0.106%	0.627%	0.004%	13076
(V)	0.085%	0.176%	0.108%	0.629%	0.001%	13074
(V)*	0.086%	0.183%	0.110%	0.618%	0.004%	12921
(VI)	0.000%	0.000%	0.503%	0.497%	0.000%	12787
(VI)*	0.000%	0.002%	0.498%	0.500%	0.000%	12787

246
247 automatic heuristic, one could use manual annotation for selected
248 comparisons [5].

250 3 Evaluating Existing Test Generators

252 Section 1 introduces potential limitations of existing input generators.
253 To evaluate this claim, I test representatives from a subset of
254 the proposed categories against the C compiler clang. I evaluate
255 the different approaches according to *input correctness* as measured
256 through manual annotation of additional instrumentation in the
257 target.

258 Table 1 presents the ratios of inputs generated by the evaluated
259 approaches that reach a certain step during program execution.
260 Approach (IV) is represented by a pure grammar-based fuzzer. For
261 (V), outputs from the grammar fuzzer are used both unchanged and
262 binary mutated. (VI) is represented by a coverage-guided grammar
263 fuzzer. All representatives were run twice, once with only self-
264 created seeds, and once with an additional set of 10 high-quality
265 manually written seeds using different aspects of the target lan-
266 guage.

267 The results confirm the suspected inability of the presented
268 approaches to test program parts past the semantic checks. They do
269 however have significant limitations and only partially adhere to
270 general evaluation best practices [17]: I am only testing against one
271 implementation, for one target requiring highly structured inputs.
272 I am evaluating only on *input correctness*, in coarse-grained steps,
273 and not against a specification. Additionally, the examples were
274 only run once for 12 hours each.

275 In future work, I want to expand on these experiments by extending
276 the number and kind of targets, measurements, and evaluation
277 robustness. I further want to test additional approaches and multiple
278 implementations in each; along with test generators purpose-built
279 for a specific target.

281 4 Reaching the Rest of the Tree

283 The results in Section 3 suggest that current approaches for test gen-
284 erators are limited — large parts of clang remain untested. Section 1
285 provides a conceptual explanation for these limitations. Building on
286 top of these, I present three approaches that may extend the parts
287 of a PUT that are effectively tested by automated test generators:

- 288 (1) **Reach down from the top:** Similarly to out-of-grammar
289 fuzzing, I want to explore out-of-language fuzzing, where

291 constraints of a full specification are iteratively and deliber-
292 ately violated, while the remaining constraints are still
293 fulfilled. This approach could further be improved by out-of-
294 grammar construction of the inputs that are evaluated for
295 constraints, giving the fuzzer theoretical reach from the top
296 of the ladder all the way to the ground.

- 297 (2) **Incremental steps towards the top:** Writing fully correct
298 specifications for tools like Fandango is considerable work
299 due to a current lack of automated specification generation.
300 However, providing a few simple constraints to an existing
301 grammar may be enough to allow a fuzzer to reach code it
302 was previously unable to test. This by design will additionally
303 produce inputs that do not match the full specification, thus
304 leading to a similar effect as approach (1).
- 305 (3) **Targeted testing:** Due to the availability of tools that able to
306 use full input specifications, I would like to explore fuzzers
307 producing inputs that reach higher interestingness in their
308 test cases by steering input generation towards boundaries
309 of both the input semantics (i.e. the constraints) and the
310 syntax (the structure).

312 5 Conclusion

313 In this work I discuss the limitations of existing test generations
314 to test programs that require highly structured inputs. Specifically,
315 I present two measurements to evaluate different test generation
316 approaches. *Input correctness* measures either how far an input
317 gets along the input processing pipeline of a target program, or to
318 what extent it fulfills analyst-given lexical, syntactic and semantic
319 constraints. *Input interestingness* measures how close an input is to
320 equivalence class boundaries.

321 I then present results of an initial experiment showing that cur-
322 rent generic test generators fail at testing significant parts of a
323 C compiler, as measured by *input correctness*. Finally, I propose
324 three approaches to creating test generators that are able to test
325 previously untestable program parts. To achieve improved *input*
326 *correctness*, systematically violate semantic and syntactic rules of
327 an input specification or extend current structure-aware test gener-
328 ators based on context-free grammars with additional, incomplete
329 semantic constraints. To produce inputs with increased *input inter-*
330 *estingness*, prioritize inputs closer to equivalence class boundaries.

332 References

- [1] [n. d.] *gcov—a Test Coverage Program*. <https://gcc.gnu.org/onlinedocs/gcc/Gcov.html>
- [2] [n. d.] *SanitizerCoverage*. <https://clang.llvm.org/docs/SanitizerCoverage.html>
- [3] Alfred V. Aho and Thomas G. Peterson. 1972. A Minimum Distance Error-Correcting Parser for Context-Free Languages. *SIAM J. Comput.* 1, 4 (1972), 305–312. arXiv:<https://doi.org/10.1137/0201022> doi:10.1137/0201022
- [4] Cornelius Aschermann, Tommaso Frassetto, Thorsten Holz, Patrick Jauernig, Ahmad-Reza Sadeghi, and Daniel Teuchert. 2019. NAUTILUS: Fishing for Deep Bugs with Grammars, In NDSS. *NDSS Symposium* 19, 337. doi:10.14722/ndss.2019.23412
- [5] Cornelius Aschermann, Sergej Schumilo, Ali Abbasi, and Thorsten Holz. 2020. Ijon: Exploring Deep State Spaces via Fuzzing, In *2020 IEEE Symposium on Security and Privacy (SP)*. 1597–1612. doi:10.1109/SP40000.2020.00117
- [6] Cornelius Aschermann, Sergej Schumilo, Tim Blazytko, Robert Gawlik, and Thorsten Holz. 2019. REDQUEEN: Fuzzing with Input-to-State Correspondence, In *Proceedings of the 2019 Network and Distributed System Security Symposium (NDSS)*, Vol. 19. 1–15.
- [7] Bachir Bendrisou, Cristian Cadar, and Alastair F. Donaldson. 2025. Grammar Mutation for Testing Input Parsers. *ACM Trans. Softw. Eng. Methodol.* 34, 4, Article 116 (April 2025), 21 pages. doi:10.1145/3708517

341
342
343
344
345
346
347
348

- 349 [8] Aarnav Bos. 2025. Autarkie. <https://github.com/R9295/autarkie/>
- 350 [9] Harrison Green, Claire Le Goues, and Fraser Brown. 2025. FrameShift: Learning
351 to Resize Fuzzer Inputs Without Breaking Them. arXiv:2507.05421 [cs.CR] <https://arxiv.org/abs/2507.05421>
- 352 [10] Valentin Huber. 2024. Differential Fuzzing on coreutils Using LibAFL. (25 06
353 2024).
- 354 [11] Valentin Huber. 2025. *FTZ: A State-Inferencing Fuzzer for the TCP/IP Stack of Zephyr*.
355 Master's thesis. Zürich University of Applied Sciences ZHAW.
- 356 [12] Romain Malmain, Andrea Fioraldi, and Aurélien Francillon. 2024. LibAFL QEMU:
357 A Library for Fuzzing-oriented Emulation. In *BAR 2024, Workshop on Binary
358 Analysis Research, colocated with NDSS 2024*. San Diego (CA), United States.
359 <https://hal.science/hal-04500872>
- 360 [13] Barton P. Miller, Lars Fredriksen, and Bryan So. 1990. An Empirical Study
361 of the Reliability of UNIX Utilities. *Commun. ACM* 33, 12 (12 1990), 32–44.
362 doi:10.1145/96267.96279
- 363 [14] Rohan Padhye, Caroline Lemieux, Koushik Sen, Mike Papadakis, and Yves
364 Le Traon. 2019. Semantic fuzzing with zest. In *Proceedings of the 28th ACM
365 SIGSOFT International Symposium on Software Testing and Analysis* (Beijing,
366 China) (ISSTA 2019). Association for Computing Machinery, New York, NY, USA,
367 329–340. doi:10.1145/3293882.3330576
- 368
- 369
- 370
- 371
- 372
- 373
- 374
- 375
- 376
- 377
- 378
- 379
- 380
- 381
- 382
- 383
- 384
- 385
- 386
- 387
- 388
- 389
- 390
- 391
- 392
- 393
- 394
- 395
- 396
- 397
- 398
- 399
- 400
- 401
- 402
- 403
- 404
- 405
- 406
- [15] Rohan Padhye, Caroline Lemieux, Koushik Sen, Laurent Simon, and Hayawardh
Vijayakumar. 2019. FuzzFactory: domain-specific fuzzing with waypoints. *Proc.
407 ACM Program. Lang.* 3, OOPSLA, Article 174 (Oct. 2019), 29 pages. doi:10.1145/
408 3360600
- [16] Sanguthevar Rajasekaran and Marius Nicolae. 2016. An Error Correcting Parser
409 for Context Free Grammars that Takes Less Than Cubic Time. In *Language and
410 Automata Theory and Applications*, Adrian-Horia Dediu, Jan Janoušek, Carlos
411 Martin-Vide, and Bianca Truthe (Eds.). Springer International Publishing, Cham,
412 533–546.
- [17] Moritz Schloegel, Nils Bars, Nico Schiller, Lukas Bernhard, Tobias Scharnowski,
413 Addison Crump, Arash Ale-Ebrahim, Nicolai Bissantz, Marius Muench, and
414 Thorsten Holz. 2024. SoK: Prudent Evaluation Practices for Fuzzing. In *2024 IEEE
415 Symposium on Security and Privacy (SP)*. 1974–1993. doi:10.1109/SP54263.2024.
416 00137
- [18] Michał Zalewski. [n. d.]. *American Fuzzy Lop - Whitepaper*. [https://lcamtuf.
418 coredump.cx/afl/technical_details.txt](https://lcamtuf.
417 coredump.cx/afl/technical_details.txt)
- [19] José Antonio Zamudio Amaya, Marius Smytzek, and Andreas Zeller. 2025. FANDANGO:
419 Evolving Language-Based Testing. *Proc. ACM Softw. Eng.* 2, ISSTA,
420 Article ISSTA040 (June 2025), 23 pages. doi:10.1145/3728915

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009