

Barely Broken – Fuzzing with Almost Valid Inputs

Valentin Huber

valentin.huber@cispa.de

CISPA Helmholtz Center for Information Security

Germany

Abstract

1.write me

CCS Concepts

- **Do Not Use This Code → Generate the Correct Terms for Your Paper;** *Generate the Correct Terms for Your Paper;* Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Keywords

Fuzzing, Grammar Fuzzing, Language Fuzzing

ACM Reference Format:

Valentin Huber. 2018. Barely Broken – Fuzzing with Almost Valid Inputs. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 4 pages. <https://doi.org/XXXXXXX.XXXXXXX>

2.change title

1 Introduction

Automated test generators such as fuzzers have been proven to be highly effective at finding software errors along these steps. However, not all test generators provide inputs that test the same parts of the system under test (SUT). [12]

Program behavior under a certain input shows similarities observed in a wide range of programs. Program input is typically provided as unstructured data (such as a binary blob containing data for an image). This input is first parsed into data structures that can then be processed by the program's business logic. For an input to be fully processed by a SUT, the linearized data in this input needs to adhere to an expected input structure, i. e. it needs to be parseable and pass all input validation checks.

Different programs have vastly different requirements for their inputs: zip is designed to process any data during compression, while compilers need to be very strict about the input they allow. Generally, one can distinguish between three steps in input validation, correlating to the examples given in Figure 1. They are increasingly less likely to be fulfilled by random data:

- (1) **Lexing:** Random bytes are unlikely to produce exclusively valid input tokens, such as valid keywords, operators, or

50 Permission to make digital or hard copies of part or all of this work for personal or
51 Unpublished working draft. Not for distribution.
52 Copyright held by author or owner(s). All rights reserved. Work partially funded by DFG, grant number 53
54 314720020. Adapted from the proceedings of the 2018 ACM SIGART Conference on Artificial Intelligence
55 and Safety, held in Woodstock, NY, USA, June 2018. This version is available at https://doi.org/XXXXXXX.XXXXXXX
56 © 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.
57 ACM ISBN 978-1-4503-XXXX-X/2018/06
58 https://doi.org/XXXXXXX.XXXXXXX

```
VBq - "7.6ark w;zu%6$K>%0TV"ryD*
&& float += % " ( ; == [
int main() { return i; }
int main() { return 0; }
```

Figure 1: Inputs for a C compiler that get rejected by the lexer, are syntactically and semantically invalid, and are fully valid, respectively.

variable names. The rules for this step can be given as a list of valid input tokens.

(2) **Syntax:** Even if one would provide a random list of exclusively valid tokens, they likely do not appear in the right structures. The rules for this step are typically given as a context-free grammar.

(3) **Semantics:** Programming languages require many additional checks, such as that each variable is defined before it is used. Recent works propose giving such additional constraints as code evaluated on elements of the grammar [20]. Various approaches have been suggested to produce inputs that adhere to some or all of these constraints [4, 6–10, 15, 20].

In this work, I present the following contributions:

- (1) I propose two alternative ways of *evaluating test generators* by examining the inputs they produce for *correctness* and *interestingness*, along with proposals for various approaches to measuring them.
- (2) I argue why I suspect that significant portions of the input space of SUTs requiring highly structured inputs remain untested, along with results from initial experiments confirming my suspicion.
- (3) Finally, I provide three suggested approaches to closing this gap.

I explore how to evaluate suspected limitations in these approaches for testing SUTs expecting highly structured inputs, along with potential solutions.

3.Talk about prior research somewhere, even if unpublished?

2 Evaluating Test Input Quality

Evaluating test generators can be done in multiple, orthogonal ways. The most common measurement is the code coverage reached when executing a certain input on an instrumented implementation of the SUT. But while code coverage is comparatively easy to measure due to the availability of instrumentation passes in popular compilers [1, 2] **citation format of these** and binary-only fuzzing tools [13], it has two fundamental limitations: (1) It assigns equal weight to

117 all branches and (2) assumes any value covering a new branch is
 118 equally interesting.

119 While code coverage is often evaluated dynamically during a
 120 fuzzing campaign to aid the fuzzer in progressing through a pro-
 121 gram, this has to be distinguished from using code coverage as a
 122 technique to evaluate fuzzer performance.

123 2.1 Input Correctness

125 Limitation (1) in evaluating a fuzzer's performance means that
 126 fuzzers that reach far into the SUT but only provide valid inputs
 127 may receive the same score as a fuzzer that tests all error paths
 128 in the initial input validation steps of a program. Alternatively, I
 129 propose evaluating fuzzers for *input correctness*. The correctness of
 130 an input with regards to a SUT can be evaluated in two ways:

- 131 (1) First, one can test inputs on an implementation of a SUT,
 measuring the ratio of accepted to rejected inputs. This can
 further be refined by annotating or instrumenting the code
 to record the step at which a certain input is rejected. This
 can be achieved with one of the following:
 - 137 (a) Manual annotations of steps along the path through the
 program, similar to [5, 16]
 - 138 (b) Selective coverage instrumentation of entire program parts
 as a very coarse measurement [15]
 - 139 (c) I further propose the following heuristic: Based on a di-
 verse set of (fuzzer-produced) inputs, *instrument the edges*
 executed by all valid inputs.
- 144 (2) Alternatively, inputs can be evaluated against an abstract
 specification of the language expected by a SUT. Such speci-
 fications are usually given in natural language, but recent
 works have explored how to express such specification in a
 structured way, as a combination of a context-free grammar
 and additional constraints over nodes in this grammar [20].
 Evaluating the correctness against such a specification may
 be done in different steps:
 - 152 (a) For inputs that cannot be parsed with the provided gram-
 mar, existing literature provides algorithm to measure how
 close to valid the input is. [3, 17]
 - 154 (b) For inputs that are parsed, the ratio of fulfilled to violated
 constraints can be used to assess the correctness of an
 input.

158 2.2 Input Interestingness

160 Not all inputs are equally interesting, even if they are all fully
 161 correct. According to the testing principle of equivalence classes,
 162 incrementing a byte is unlikely to trigger a bug. This is of course
 163 unless this new byte is at the edge of legality for the SUT, for
 164 example by changing this byte so that it is no longer in a valid
 165 range. This is known as boundary testing: If a requirement allows
 166 values between 5 and 100, the likelihood of discovering a problem
 167 is higher with values 4, 5, 6, 99, 100, and 101 compared to values
 168 between and outside. Special values such as 0, -1, or the maximum
 169 possible number may further be interesting.

170 Previously, it was non-trivial to evaluate how close an input part
 171 is to such a constraint boundary. With the availability of formal
 172 and machine-readable specifications, such as described above, this
 173 now becomes possible. An input (part) is more interesting, the

175 closer to the boundaries provided by the specification it is. This
 176 can be applied to semantic constraints, but also syntactic structure:
 177 If a node can be repeated an arbitrary number of times, testing
 178 0 repetitions, 1 repetition and a large number of repetitions is
 179 interesting. This can be reinterpreted as syntactic boundary testing.

181 3 Evaluating Existing Test Generators

182 I do not know of any work systematically evaluating correctness
 183 and interestingness of the inputs produced by different test genera-
 184 tors. Considering that previous work found significant problems
 185 with the evaluation of fuzzers [18], evaluating different approaches
 186 along the axes outlined in Section 2 may lead to greater insight into
 187 what approaches are actually evaluating interesting parts of the
 188 target; or more importantly, which approaches produce individual
 189 in parts of the input space no other test generator reaches.

190 For an initial experiment, I ran the grammar fuzzer Nautilus¹ [4]
 191 against the JavaScript runtime QuickJS and the C compiler clang
 192 with respective grammars. According to the first correctness metric
 193 outlined above, I recorded the correctness of inputs produced by
 194 the fuzzer.

195 Running against QuickJS, Nautilus produced approximately 45%
 196 inputs that could not be parsed, 5% inputs that resulted in an ex-
 197 ception, and 50% inputs that were evaluated without an error. For
 198 clang, the results show considerable differences: Only about 50%
 199 of inputs made it past the lexing step, and not a single one made it
 200 past the parser. One can therefore see how the compiler, assembler,
 201 and linker step of clang are never tested with this fuzzing setup.
 202 This is likely because clang uses custom lexing and parsing logic
 203 that integrates semantic checks. So while the inputs produced by
 204 Nautilus are discarded by clang, they by construction can be parsed
 205 by the abstract grammar.

206 The differences in complexity of constraints between JavaScript
 207 and C do not explain the difference in result. It seems much more
 208 likely that these differences are because of the permissiveness and
 209 dynamic typing of JavaScript, while many deeper parts of the inter-
 210 preter are not actually getting tested. Here, evaluating test correct-
 211 ness against a specification may be more appropriate.

212 For clang, Nautilus does not seem to be able to generate test
 213 inputs that fully compile and thus test all program parts; other
 214 approaches should be used, at least in addition to Nautilus.

215 5.Add plot of correctness levels against time

217 4 The Ladder to Success: Categorizing 218 Approaches

219 One can categorize most fuzzers along the correctness of the inputs
 220 they provide. Imagine the distribution of bugs like the distribution
 221 of berries in a bush, with berries higher up in the bush representing
 222 more correct inputs, as shown in Figure 2. Historically, a set of
 223 fundamentally different approaches was proposed. These produce
 224 inputs that have varying distributions of input correctness and
 225 interestingness and therefore test different parts of the program
 226 and check for different kinds of software errors. The following
 227 section provides a categorization of approaches along with the
 228 program parts they test, as indicated in Figure 2. The exact extent

229 ¹Nautilus version 2.0 through libafl_nautilus; which notably no longer provides out-of-grammar fuzzing

of the figurative reach of each fuzzer remains unknown; this figure is not to scale, including the amount of overlap.

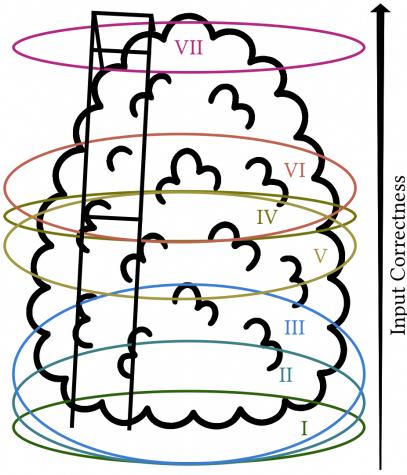


Figure 2: Bush representing the input space of a program, where increasing height represents increasing correctness of a certain input

The initial approach was using fully random data [14]. As discussed above, this will only test the outermost input validation, equivalent to the very low-hanging fruit (as represented by the ellipsis labeled I).

The next major step was the introduction of naïve mutational coverage-guided fuzzing (CGF) in American Fuzzy Lop (AFL) [19]. By reaching more and more coverage, these fuzzers "learn" the structure the inputs they produce require to reach more coverage. They are however severely limited by the kind of structure they can produce: Even multi-byte comparison operations requires them to guess the value. These fuzzers reached more of the low-hanging fruit in the tree of bugs with no help in the form of an input specification from the analyst (II).

Some of these limitations were mitigated to a point by extending CGF with additional instrumentation. By logging the values input parts are compared to, the fuzzer can mutate the input again to fix these parts and therefore pass the checks [6]. This further allows CGF to pass simple checksums. More recent advancements allow CGF to automatically fix length fields based on sudden coverage drop-offs [10]. These improvements allow the fuzzer to produce more correct inputs and reach further up the tree of bugs (III).

However, certain structure still cannot be learned by fuzzers from feedback, which is why [6.cite](#) grammar fuzzers were introduced [4]. These start from a different place, producing inputs that pass at least the lexing and parsing steps of a SUT². By providing the grammar, the analyst "helps" the fuzzer onto the first step of a ladder on the bush where it reaches left and right (IV).

From this next step on the ladder, fuzzers can reach further down by producing inputs that do not (fully) adhere to the grammar, by mutating the grammar or the trees after production (V) [4, 7]. There have been initial steps towards extending grammar-based fuzzers

²At least in principle; unless constraint validation is intermixed with parsing

to reach further up, for example by using the same instrumentation on value comparisons as described above (VI) [8].

Language-based fuzzing extends context-free grammars with additional constraints that ensure the inputs produced are not only syntactically but also semantically correct. This correlates to an analyst providing all information necessary for the fuzzer to produce fully correct input (assuming a complete and correct specification), which is equivalent to the topmost step of the ladder (VII).

Fuzzers based on symbolic execution can be thought of as systematically walking each branch of the bush. In some of my previous work, I have discussed the limitations of fuzzers relying on symbolic execution [11].

There further exists a list of fuzzers purpose-built for a certain target specification. They contain information about the expected input format, encoded in their logic. And while they may be highly optimized, they are fundamentally equivalent to a generic fuzzer given the same information in the form of grammar or language specifications and should be categorized accordingly.

5 Reaching the Rest of the Tree

The combined input space coverage as seen in Figure 2 seems to suggest that existing approaches may not reach all program logic; although this will have to be validated with a survey based on the measurements as described in Section 2. The following approaches may do so:

- (1) **Reach down from the top:** Similarly to out-of-grammar fuzzing, I want to explore out-of-language fuzzing, where constraints of a full specification are iteratively and deliberately violated, while the remaining constraints are still fulfilled. This approach could further be improved by out-of-grammar construction of the inputs that are evaluated for constraints.
- (2) **Incremental steps towards the top:** Writing fully correct specifications for tools like Fandango is considerable work due to a current lack of automated specification generation. However, providing a few simple constraints to an existing grammar may be enough to allow a fuzzer to reach code it was previously unable to test. This by design will additionally produce inputs that do not match the full specification, thus leading to a similar effect as approach (1).
- (3) **Targeted testing:** Due to the availability of tools able to use full input specifications, I would like to explore fuzzers producing inputs that reach higher interestingness in their test cases by steering input generation towards boundaries of both the input semantics (i.e. the constraints) and the syntax (the structure).

6 Conclusion

In this work, I outline how suspected limitations of existing test generators; namely their inability to test parts of target software that requires highly structured inputs to be reached. Initial results on a JavaScript interpreter and a C compiler suggest significant parts of such targets may not be tested at all by current test generators.

I further propose multiple measurements to evaluate the inputs of existing and future test generators with regards to their correctness against an implementation and an abstract specification, along

349 with their interestingness as measured by distance to boundaries
 350 of equivalence classes.

351 Finally, I propose multiple approaches of how to design future
 352 test generators to fill gaps in existing work, agnostic to specification
 353 and implementation.

354 References

- [1] [n. d.]. *gcov—a Test Coverage Program*. <https://gcc.gnu.org/onlinedocs/gcc/Gcov.html>
- [2] [n. d.]. *SanitizerCoverage*. <https://clang.llvm.org/docs/SanitizerCoverage.html>
- [3] Alfred V. Aho and Thomas G. Peterson. 1972. A Minimum Distance Error-Correcting Parser for Context-Free Languages. *SIAM J. Comput.* 1, 4 (1972), 305–312. arXiv:<https://doi.org/10.1137/0201022> doi:10.1137/0201022
- [4] Cornelius Aschermann, Tommaso Frassetto, Thorsten Holz, Patrick Jauernig, Ahmad-Reza Sadeghi, and Daniel Teuchert. 2019. NAUTILUS: Fishing for Deep Bugs with Grammars, In NDSS. *NDSS Symposium* 19, 337. doi:10.14722/ndss.2019.23412
- [5] Cornelius Aschermann, Sergej Schumilo, Ali Abbasi, and Thorsten Holz. 2020. Ijon: Exploring Deep State Spaces via Fuzzing. In *2020 IEEE Symposium on Security and Privacy (SP)*, 1597–1612. doi:10.1109/SP40000.2020.00117
- [6] Cornelius Aschermann, Sergej Schumilo, Tim Blazytko, Robert Gawlik, and Thorsten Holz. 2019. REDQUEEN: Fuzzing with Input-to-State Correspondence. In *Proceedings of the 2019 Network and Distributed System Security Symposium (NDSS)*, Vol. 19, 1–15.
- [7] Bachir Bendrissou, Cristian Cadar, and Alastair F. Donaldson. 2025. Grammar Mutation for Testing Input Parsers. *ACM Trans. Softw. Eng. Methodol.* 34, 4, Article 116 (April 2025), 21 pages. doi:10.1145/3708517
- [8] Aarnav Bos. 2025. Autarkie. <https://github.com/R9295/autarkie/>
- [9] Andrea Fioraldi, Dominik Maier, Heiko Eißfeldt, and Marc Heuse. 2020. AFL++: Combining Incremental Steps of Fuzzing Research. In *14th USENIX Workshop on Offensive Technologies (WOOT 20)*. USENIX Association.
- [10] Harrison Green, Claire Le Goues, and Fraser Brown. 2025. FrameShift: Learning to Resize Fuzzer Inputs Without Breaking Them. arXiv:2507.05421 [cs.CR] <https://arxiv.org/abs/2507.05421>
- [11] Valentin Huber. 2023. Challenges and Mitigation Strategies in Symbolic Execution Based Fuzzing Through the Lens of Survey Papers. (15 12 2023). <https://github.com/riesentoaster/review-symbolic-execution-in-fuzzing/releases/download/v1.0/Huber-Valentin-Challenges-and-Mitigation-Strategies-in-Symbolic-Execution-Based-Fuzzing-Through-the-Lens-of-Survey-Papers.pdf>
- [12] Sanoop Mallissery and Yu-Sung Wu. 2023. Demystify the Fuzzing Methods: A Comprehensive Survey. *ACM Comput. Surv.* 56, 3, Article 71 (Oct. 2023), 38 pages. doi:10.1145/3623375
- [13] Romain Malmain, Andrea Fioraldi, and Aurélien Francillon. 2024. LibAFL QEMU: A Library for Fuzzing-oriented Emulation. In *BAR 2024, Workshop on Binary Analysis Research, colocated with NDSS 2024*. San Diego (CA), United States. <https://hal.science/hal-04500872>
- [14] Barton P. Miller, Lars Fredriksen, and Bryan So. 1990. An Empirical Study of the Reliability of UNIX Utilities. *Commun. ACM* 33, 12 (12 1990), 32–44. doi:10.1145/96267.96279
- [15] Rohan Padhye, Caroline Lemieux, Koushik Sen, Mike Papadakis, and Yves Le Traon. 2019. Semantic fuzzing with zest. In *Proceedings of the 28th ACM SIGSOFT International Symposium on Software Testing and Analysis* (Beijing, China) (ISSTA 2019). Association for Computing Machinery, New York, NY, USA, 329–340. doi:10.1145/3293882.3330576
- [16] Rohan Padhye, Caroline Lemieux, Koushik Sen, Laurent Simon, and Hayawardh Vijayakumar. 2019. FuzzFactory: domain-specific fuzzing with waypoints. *Proc. ACM Program. Lang.* 3, OOPSLA, Article 174 (Oct. 2019), 29 pages. doi:10.1145/3360600
- [17] Sanguthevar Rajasekaran and Marius Nicolae. 2016. An Error Correcting Parser for Context Free Grammars that Takes Less Than Cubic Time. In *Language and Automata Theory and Applications*, Adrian-Horia Dediu, Jan Janoušek, Carlos Martin-Vide, and Bianca Truthe (Eds.). Springer International Publishing, Cham, 533–546.
- [18] Moritz Schloegl, Nils Bars, Nico Schiller, Lukas Bernhard, Tobias Scharnowski, Addison Crump, Arash Ale-Elbrahim, Nicolai Bissantz, Marius Muench, and Thorsten Holz. 2024. SoK: Prudent Evaluation Practices for Fuzzing. In *2024 IEEE Symposium on Security and Privacy (SP)*. 1974–1993. doi:10.1109/SP54263.2024.00137
- [19] Michał Zalewski. [n. d.]. *American Fuzzy Lop - Whitepaper*. https://lcamtuf.coredump.cx/afl/technical_details.txt
- [20] José Antonio Zamudio Amaya, Marius Smytzek, and Andreas Zeller. 2025. FANDANGO: Evolving Language-Based Testing. *Proc. ACM Softw. Eng.* 2, ISSTA, Article ISSTA040 (June 2025), 23 pages. doi:10.1145/3728915

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009